

Production-Optimized Dockerfiles

These Dockerfiles are optimized for production deployment with security, performance, and size considerations.

Directory Structure

Create a `docker/` directory in your project root:

```
threatcompass/
├── docker/
│   ├── Dockerfile.flask-app
│   ├── Dockerfile.celery-worker
│   ├── Dockerfile.celery-beat
│   └── requirements/
│       ├── base.txt
│       ├── production.txt
│       └── development.txt
├── requirements.txt -> docker/requirements/production.txt
└── requirements-dev.txt -> docker/requirements/development.txt
```

Flask Application Dockerfile

dockerfile

docker/Dockerfile.flask-app

Production-optimized Flask application container

Use Python 3.11 slim image for smaller size

FROM python:3.11-slim **as** base

Set environment variables

ENV PYTHONUNBUFFERED=1 \
PYTHONDONTWRITEBYTECODE=1 \
PIP_NO_CACHE_DIR=1 \
PIP_DISABLE_PIP_VERSION_CHECK=1 \
PIP_DEFAULT_TIMEOUT=100

Install system dependencies

RUN apt-get update && apt-get install -y --no-install-recommends \

Build dependencies

gcc \

g++ \

PostgreSQL client library

libpq-dev \

For health checks

curl \

For git dependencies (if any)

git \

Clean up

&& rm -rf /var/lib/apt/lists/* \

&& apt-get clean

```
# =====:
# Builder stage - Install Python dependencies
# =====:
FROM base as builder

# Create virtual environment
RUN python -m venv /opt/venv
ENV PATH="/opt/venv/bin:$PATH"

# Copy requirements first for better caching
COPY docker/requirements/production.txt /tmp/requirements.txt

# Install Python dependencies
RUN pip install --upgrade pip setuptools wheel && \
    pip install -r /tmp/requirements.txt

# =====:
# Production stage - Final image
# =====:
FROM base as production

# Create non-root user for security
RUN groupadd --gid 1000 appgroup && \
    useradd --uid 1000 --gid appgroup --shell /bin/bash --create-home appuser

# Copy virtual environment from builder stage
COPY --from=builder /opt/venv /opt/venv
ENV PATH="/opt/venv/bin:$PATH"
```

Set work directory

WORKDIR /app

Copy application code

COPY --chown=appuser:appgroup . .

Remove unnecessary files for production

RUN rm -rf tests/ \
docker/requirements/development.txt \
requirements-dev.txt \
.git/ \
.pytest_cache/ \
__pycache__/ \
*.pyc \
.coverage

Create directories for logs and temporary files

RUN mkdir -p /app/logs /app/tmp && \
chown -R appuser:appgroup /app

Switch to non-root user

USER appuser

Health check

HEALTHCHECK --interval=30s --timeout=10s --start-period=30s --retries=3 \
CMD curl -f http://localhost:5000/health || exit 1

Expose port

EXPOSE 5000

Default command - use gunicorn for production

```
CMD ["gunicorn", \  
    "--bind", "0.0.0.0:5000", \  
    "--workers", "4", \  
    "--worker-class", "gevent", \  
    "--worker-connections", "1000", \  
    "--timeout", "120", \  
    "--keepalive", "5", \  
    "--max-requests", "1000", \  
    "--max-requests-jitter", "100", \  
    "--preload", \  
    "--access-logfile", "-", \  
    "--error-logfile", "-", \  
    "--log-level", "info", \  
    "app:app"]
```

Celery Worker Dockerfile

dockerfile

docker/Dockerfile.celery-worker

Production-optimized Celery worker container

FROM python:3.11-slim **as** base

ENV PYTHONUNBUFFERED=1 \
PYTHONDONTWRITEBYTECODE=1 \
PIP_NO_CACHE_DIR=1 \
PIP_DISABLE_PIP_VERSION_CHECK=1 \
PIP_DEFAULT_TIMEOUT=100 \
C_FORCE_ROOT=1

Install system dependencies

RUN apt-get update && apt-get install -y --no-install-recommends \
gcc \
g++ \
libpq-dev \
curl \
&& rm -rf /var/lib/apt/lists/* \
&& apt-get clean

=====

Builder stage

=====

FROM base **as** builder

RUN python -m venv /opt/venv

ENV PATH="/opt/venv/bin:\$PATH"

COPY docker/requirements/production.txt /tmp/requirements.txt

RUN pip install --upgrade pip setuptools wheel && \
pip install -r /tmp/requirements.txt

=====:

Production stage

=====:

FROM base **as** production

Create non-root user

RUN groupadd --gid 1000 celery && \
useradd --uid 1000 --gid celery --shell /bin/bash --create-home celery

Copy virtual environment

COPY --from=builder /opt/venv /opt/venv

ENV PATH="/opt/venv/bin:\$PATH"

WORKDIR /app

Copy application code

COPY --chown=celery:celery ..

Clean up unnecessary files

RUN rm -rf tests/ \
docker/requirements/development.txt \
requirements-dev.txt \
.git/ \
.pytest_cache/ \

__pycache__/\

*.pyc \

.coverage

Create log directory

RUN mkdir -p /app/logs && chown -R celery:celery /app

USER celery

Health check for Celery worker

HEALTHCHECK --interval=60s --timeout=30s --start-period=60s --retries=3 \

CMD celery -A app.celery inspect ping -d celery@\$HOSTNAME || exit 1

Default command

CMD ["celery", \

"-A", "app.celery", \

"worker", \

"--loglevel=info", \

"--concurrency=4", \

"--prefetch-multiplier=1", \

"--max-tasks-per-child=1000", \

"--time-limit=7200", \

"--soft-time-limit=3600"]

Celery Beat Dockerfile

dockerfile

```
# docker/Dockerfile.celery-beat  
# Production-optimized Celery beat scheduler container
```

```
FROM python:3.11-slim as base
```

```
ENV PYTHONUNBUFFERED=1 \  
    PYTHONDONTWRITEBYTECODE=1 \  
    PIP_NO_CACHE_DIR=1 \  
    PIP_DISABLE_PIP_VERSION_CHECK=1 \  
    PIP_DEFAULT_TIMEOUT=100 \  
    C_FORCE_ROOT=1
```

```
# Install system dependencies
```

```
RUN apt-get update && apt-get install -y --no-install-recommends \  
    gcc \  
    g++ \  
    libpq-dev \  
    && rm -rf /var/lib/apt/lists/* \  
    && apt-get clean
```

```
# =====
```

```
# Builder stage
```

```
# =====
```

```
FROM base as builder
```

```
RUN python -m venv /opt/venv
```

```
ENV PATH="/opt/venv/bin:$PATH"
```

COPY docker/requirements/production.txt /tmp/requirements.txt

RUN pip install --upgrade pip setuptools wheel && \
pip install -r /tmp/requirements.txt

=====:

Production stage

=====:

FROM base **as** production

Create non-root user

RUN groupadd --gid 1000 celery && \
useradd --uid 1000 --gid celery --shell /bin/bash --create-home celery

Copy virtual environment

COPY --from=builder /opt/venv /opt/venv

ENV PATH="/opt/venv/bin:\$PATH"

WORKDIR /app

Copy application code

COPY --chown=celery:celery ..

Clean up unnecessary files

RUN rm -rf tests/ \
docker/requirements/development.txt \
requirements-dev.txt \
.git/ \
.pytest_cache/ \
__pycache__/

*.pyc \

.coverage

Create directories

RUN mkdir -p /app/logs /app/celerybeat-schedule && \

chown -R celery:celery /app

USER celery

Health check - check if beat process is running

HEALTHCHECK --interval=60s --timeout=30s --start-period=60s --retries=3 \

CMD pgrep -f "celery.*beat" || exit 1

Default command

CMD ["celery", \

"-A", "app.celery", \

"beat", \

"--loglevel=info", \

"--schedule=/app/celerybeat-schedule/celerybeat-schedule", \

"--pidfile=/app/celerybeat-schedule/celerybeat.pid"]

Requirements Files

Base Requirements (`docker/requirements/base.txt`)

txt

Core Flask dependencies

Flask==2.3.3

Flask-SQLAlchemy==3.0.5

Flask-Login==0.6.3

Flask-WTF==1.1.1

Flask-Migrate==4.0.5

Flask-Limiter==3.5.0

Flask-CORS==4.0.0

Flask-Caching==2.1.0

Database

SQLAlchemy==2.0.21

psycopg2-binary==2.9.7

alembic==1.12.0

Celery and Redis

celery==5.3.2

redis==5.0.1

Security

cryptography==41.0.4

bcrypt==4.0.1

HTTP and API

requests==2.31.0

urllib3==2.0.5

Data processing

pandas==2.1.1
numpy==1.25.2

Validation
marshmallow==3.20.1
WTForms==3.0.1

AWS SDK
boto3==1.28.57
botocore==1.31.57

Utilities
python-dotenv==1.0.0
click==8.1.7

Production Requirements ([docker/requirements/production.txt](#))

txt

-r base.txt

Production WSGI server

gunicorn[gevent]==21.2.0

gevent==23.7.0

Production monitoring

sentry-sdk[flask]==1.32.0

Performance

greenlet==2.0.2

Email

flask-mail==0.9.1

Development Requirements

((docker/requirements/development.txt))

txt

-r production.txt

Testing

pytest==7.4.2

pytest-cov==4.1.0

pytest-flask==1.2.0

pytest-mock==3.11.1

Code quality

black==23.9.1

flake8==6.1.0

isort==5.12.0

mypy==1.5.1

Security testing

bandit==1.7.5

safety==2.3.4

Development tools

ipython==8.15.0

flask-debugtoolbar==0.13.1

Build Scripts

Build Script ([**scripts/build-images.sh**](#)**)**

bash

```
#!/bin/bash
```

```
# scripts/build-images.sh
```

```
# Build all Docker images for ThreatCompass
```

```
set -e
```

```
# Configuration
```

```
REGISTRY=${1:-"local"}
```

```
TAG=${2:-"latest"}
```

```
REGION=${3:-"us-east-1"}
```

```
echo "Building ThreatCompass Docker images..."
```

```
echo "Registry: $REGISTRY"
```

```
echo "Tag: $TAG"
```

```
# Get Git commit SHA for tagging
```

```
GIT_SHA=$(git rev-parse --short HEAD)
```

```
# Build Flask App
```

```
echo "Building Flask App..."
```

```
docker build \
```

```
-f docker/Dockerfile.flask-app \
```

```
-t threatcompass-flask-app:$TAG \
```

```
-t threatcompass-flask-app:$GIT_SHA \
```

```
--build-arg BUILD_DATE=$(date -u +'%Y-%m-%dT%H:%M:%SZ') \
```

```
--build-arg GIT_SHA=$GIT_SHA \
```

```
.
```

Build Celery Worker

echo "Building Celery Worker..."

```
docker build \
  -f docker/Dockerfile.celery-worker \
  -t threatcompass-celery-worker:$TAG \
  -t threatcompass-celery-worker:$GIT_SHA \
  --build-arg BUILD_DATE=$(date -u +'%Y-%m-%dT%H:%M:%SZ') \
  --build-arg GIT_SHA=$GIT_SHA \
  .
```

Build Celery Beat

echo "Building Celery Beat..."

```
docker build \
  -f docker/Dockerfile.celery-beat \
  -t threatcompass-celery-beat:$TAG \
  -t threatcompass-celery-beat:$GIT_SHA \
  --build-arg BUILD_DATE=$(date -u +'%Y-%m-%dT%H:%M:%SZ') \
  --build-arg GIT_SHA=$GIT_SHA \
  .
```

echo "Build completed successfully!"

Optional: Push to ECR

```
if [ "$REGISTRY" != "local" ]; then
```

```
  echo "Pushing images to ECR..."
```

Login to ECR

```
aws ecr get-login-password --region $REGION | \
```

```
  docker login --username AWS --password-stdin $REGISTRY
```

Tag and push images

for service in flask-app celery-worker celery-beat; do

docker tag threatcompass-\$service:\$TAG \$REGISTRY/threatcompass-production

docker tag threatcompass-\$service:\$GIT_SHA \$REGISTRY/threatcompass-pr

docker push \$REGISTRY/threatcompass-production/\$service:\$TAG

docker push \$REGISTRY/threatcompass-production/\$service:\$GIT_SHA

done

echo "Images pushed to ECR successfully!"

fi

Development Build Script (`scripts/build-dev.sh`)

```
bash
```

```
#!/bin/bash
```

```
# scripts/build-dev.sh
```

```
# Build development images with development dependencies
```

```
set -e
```

```
echo "Building ThreatCompass development images..."
```

```
# Build development Flask app with hot reload
```

```
docker build \
```

```
-f docker/Dockerfile.flask-app \
```

```
--target base \
```

```
-t threatcompass-flask-app:dev \
```

```
--build-arg REQUIREMENTS_FILE=docker/requirements/development.txt \
```

```
.
```

```
echo "Development build completed!"
```

Docker Compose for Local Development

yaml

docker-compose.dev.yml

version: '3.8'

services:

web:

build:

context: .

dockerfile: docker/Dockerfile.flask-app

target: base

volumes:

- ./app
- /app/docker/requirements/

ports:

- "5000:5000"

environment:

- FLASK_ENV=development
- FLASK_DEBUG=1
- DATABASE_URL=postgresql://postgres:postgres@db:5432/threatcompass_db
- REDIS_URL=redis://redis:6379/0

depends_on:

- db
- redis

command: ["flask", "run", "--host=0.0.0.0", "--port=5000", "--reload"]

celery_worker:

build:

context: .

dockerfile: docker/Dockerfile.celery-worker

target: base

volumes:

- ./app

environment:

- DATABASE_URL=postgresql://postgres:postgres@db:5432/threatcompass_db
- REDIS_URL=redis://redis:6379/0

depends_on:

- db
- redis

command: ["celery", "-A", "app.celery", "worker", "--loglevel=debug", "--reload"]

celery_beat:

build:

context: .

dockerfile: docker/Dockerfile.celery-beat

target: base

volumes:

- ./app

environment:

- DATABASE_URL=postgresql://postgres:postgres@db:5432/threatcompass_db
- REDIS_URL=redis://redis:6379/0

depends_on:

- db
- redis

command: ["celery", "-A", "app.celery", "beat", "--loglevel=debug"]

db:

image: postgres:15

environment:

- POSTGRES_DB=threatcompass_dev
- POSTGRES_USER=postgres
- POSTGRES_PASSWORD=postgres

ports:

- "5432:5432"

volumes:

- postgres_data:/var/lib/postgresql/data

redis:

image: redis:7-alpine

ports:

- "6379:6379"

volumes:

postgres_data:

Security Optimizations

1. Distroless Images (Advanced)

dockerfile

Alternative base using distroless for maximum security

FROM gcr.io/distroless/python3-debian11

COPY --from=builder /opt/venv /opt/venv

COPY --chown=nonroot:nonroot ./app

WORKDIR /app

USER nonroot

ENV PATH="/opt/venv/bin:\$PATH"

CMD ["/opt/venv/bin/gunicorn", "app:app"]

2. Multi-stage Builds with Vulnerability Scanning

dockerfile




Add vulnerability scanning labels

LABEL org.opencontainers.image.source="https://github.com/your-org/threatcompass" 

LABEL org.opencontainers.image.description="ThreatCompass Flask Application" 

LABEL org.opencontainers.image.licenses="MIT" 

These production-optimized Dockerfiles provide:

-  **Multi-stage builds** for smaller final images
-  **Non-root users** for security
-  **Health checks** for container orchestration

- ✓ **Optimized caching** with proper layer ordering
- ✓ **Clean environments** with unnecessary files removed
- ✓ **Production-ready configurations** for Gunicorn and Celery

The CI/CD pipeline will use these Dockerfiles to build secure, efficient containers for your AWS ECS deployment.